

Northeastern University - Seattle

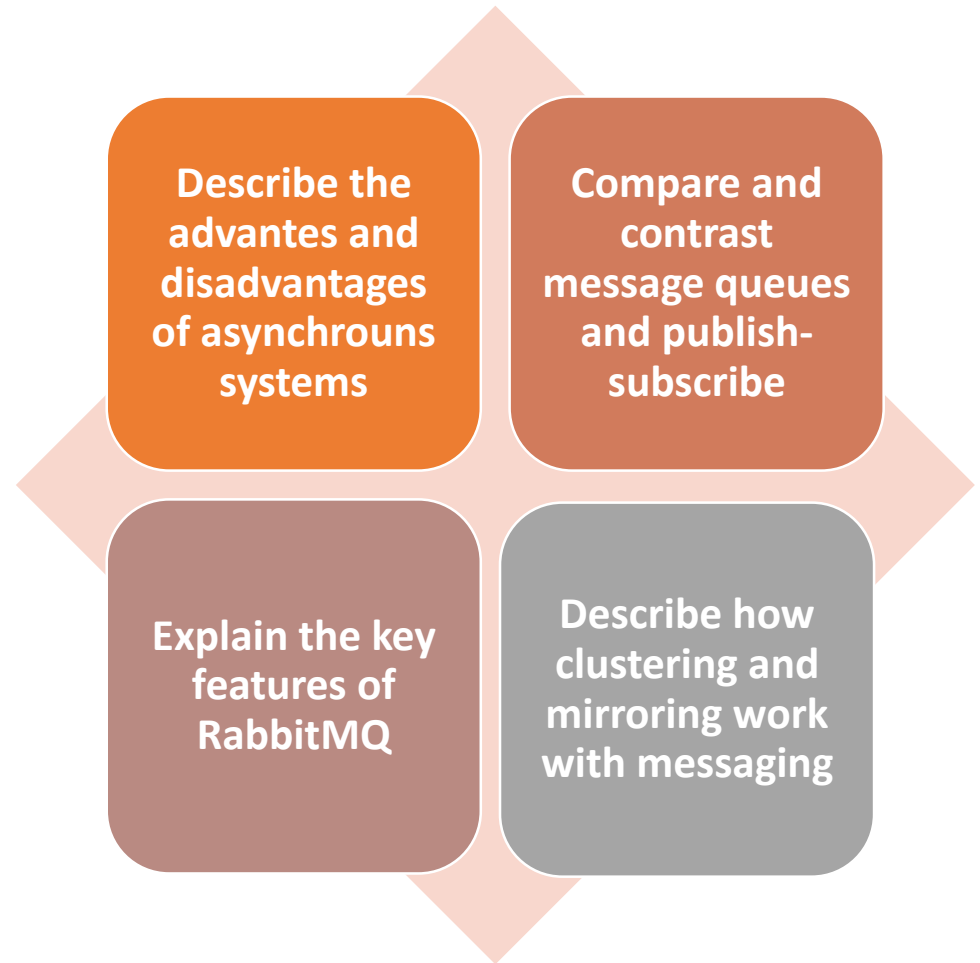


CS6650 Building Scalable Distributed Systems
Professor Ian Gorton

Building Scalable Distributed Systems

Week 8 – Asynchronous Systems

Learning Objectives



Outline

- Asynchronous systems basics
- Example: RabbitMQ
- Messaging patterns
- Clustering and mirroring

Asynchronous Systems Basics

Background

- So far we've examined communications such as:
 - Sockets
 - RPC
 - HTTP
- These work great!
 - Simple to program, eg [Thrift](#), [protocol buffers](#)
 - Fast comms
- But lead to tight coupling:
 - Client needs to know who server is
 - What happens if the server is not available?

Middleware


- Middleware is the plumbing or wiring of many applications
- Provides applications with fundamental services for distributed computing
- Insulates applications from underlying platform (OS, DBMS, etc) APIs
- Lots of middleware exists
 - Different purposes
 - Different vendors
 - Different standards and proprietary technologies
- Focus of today – Message Oriented Middleware (MOM)

Messaging - MOM

- Message Oriented Middleware (MOM) provides:
 - Asynchronous communications between processes, applications and systems
 - Send-and-forget - Delivering messages despite failures
 - Transactional Messaging - Deliver all messages in a transaction, or none
 - Persistence - Messages can be logged at the server and hence survive server failure




Messaging Primitives




- Send (queue, message)
 - Put message onto queue
- Receive (queue, message)
 - Get message from queue
- No dependency on state of receiving application on message send

Messaging Primitives: Persistence



- Receipt of message at queue implies message is written to disk log
- Removal of message from queue deletes message from disk log
- Trade-off performance versus reliability


Messaging Primitives: Transactions



```
Begin transaction
...
update database record
put message on queue
...
commit transaction
```


```
Begin transaction
...
get message from queue
update database record
...
commit transaction
```

Messaging Primitives: Transactions



- Sender and receiver do **not** share a transaction
 - Rollback on receiver does not affect the sender
- ‘Synchronous’ operations are not atomic
(Request/response is 3 transactions not 1)
 - Put to request queue
 - Get from request queue, put to response queue
 - Get from response queue

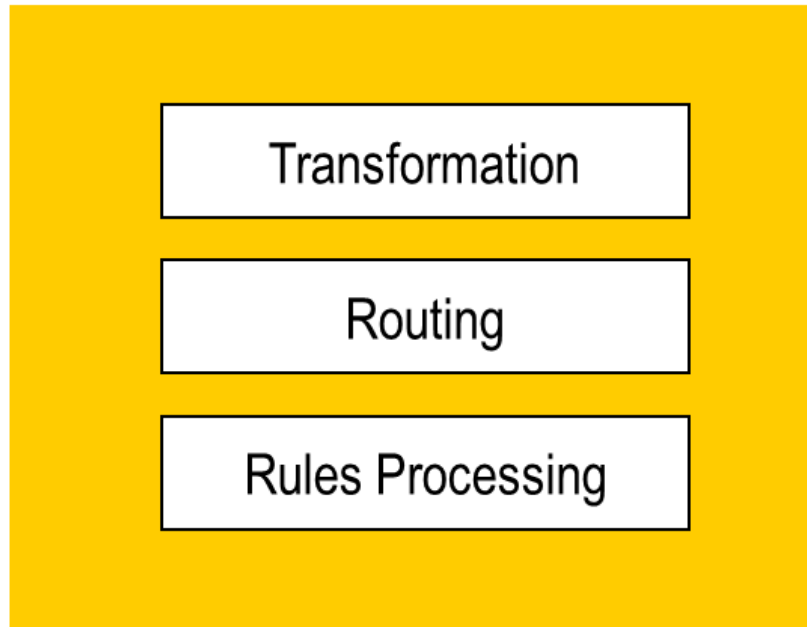

Messaging Primitives: Publish-Subscribe




- Provide 1-to-N, N-to-1, and N-to-N communications
- Messages are 'published' to logical *subjects* or *topics*
- Subscribers receive all messages from subjects they subscribe to

Hub and Spoke Architecture

Input Messages



Output Messages



Message Brokers

(A big step towards an Enterprise Service Bus (ESB))

Messaging Implementations

- Many! Many!! Many!!!
- Each has different:
 - Features
 - Strengths
 - Weaknesses
- In general there are two extreme choices:
 - Smart endpoints/dumb pipes
 - Dumb endpoints/smart pipes
- Which scales best?


Example: RabbitMQ

- Widely used, open source broker that implements AMQP protocol
- messages are published to exchanges
- Exchanges distribute messages to queues using rules called bindings.
- the broker can deliver messages to consumers subscribed to queues
- Optionally consumers fetch/pull messages from queues on demand



Examples: RabbitMQ

"Hello, world" example routing



Producer

```
import com.rabbitmq.client.Channel;
import com.rabbitmq.client.Connection;
import com.rabbitmq.client.ConnectionFactory;

public class Send {

    private final static String QUEUE_NAME = "hello";

    public static void main(String[] argv) throws Exception {
        ConnectionFactory factory = new ConnectionFactory();
        factory.setHost("localhost");
        try (Connection connection = factory.newConnection();
            Channel channel = connection.createChannel()) {
            channel.queueDeclare(QUEUE_NAME, false, false, false, null);
            String message = "Hello World!";
            channel.basicPublish("", QUEUE_NAME, null, message.getBytes("UTF-8"));
            System.out.println(" [x] Sent '" + message + "'");
        }
    }
}
```

Consumer

```
import com.rabbitmq.client.Channel;
import com.rabbitmq.client.Connection;
import com.rabbitmq.client.ConnectionFactory;
import com.rabbitmq.client.DeliverCallback;

public class Recv {


    private final static String QUEUE_NAME = "hello";

    public static void main(String[] argv) throws Exception {
        ConnectionFactory factory = new ConnectionFactory();
        factory.setHost("localhost");
        Connection connection = factory.newConnection();
        Channel channel = connection.createChannel();

        channel.queueDeclare(QUEUE_NAME, false, false, false, null);
        System.out.println(" [*] Waiting for messages. To exit press CTRL+C");

        DeliverCallback deliverCallback = (consumerTag, delivery) -> {
            String message = new String(delivery.getBody(), "UTF-8");
            System.out.println(" [x] Received '" + message + "'");
        };
        channel.basicConsume(QUEUE_NAME, true, deliverCallback, consumerTag -> { });
    }
}
```

Multiple Consumers



- Messages distributed in a round robin fashion to consumers
- Each message sent to exactly one consumer
- Consumers send message ack when processing complete
- Unack-ed messages redistributed when a consumer dies


Consumer – Manual Acknowledge

```
channel.basicQos(1); // accept only one unack-ed message at a time
```

```
DeliverCallback deliverCallback = (consumerTag, delivery) -> {  
    String message = new String(delivery.getBody(), "UTF-8");
```

```
    System.out.println(" [x] Received '" + message + "'");  
    try {  
        ProcessMessage(message);  
    } finally {  
        System.out.println(" [x] Done");  
        channel.basicAck(delivery.getEnvelope().getDeliveryTag(), false);  
    }  
};  
boolean autoAck = false;  
channel.basicConsume(TASK_QUEUE_NAME, autoAck, deliverCallback,  
consumerTag -> { });
```


Message Distribution



- Limits number of messages sent to a single consumer
- If consumer is busy, messages are sent to next consumer with capacity

Persistent/Durable Messaging

- If RabbitMQ broker fails, it loses:
 - All queues on the server
 - All messages in the queues
- For safety, we need to persist the broker configuration and messages



Persistent/Durable Messaging

```
// in producer and consumer  
boolean durable = true;  
channel.queueDeclare("task_queue", durable, false, false, null);
```

```
// and ....  
import com.rabbitmq.client.MessageProperties;
```

```
channel.basicPublish("", "task_queue",  
    MessageProperties.PERSISTENT_TEXT_PLAIN,  
    message.getBytes());
```



Publish-Subscribe

- Every message sent to all subscribers
- Utilize RabbitMQ exchanges
- Producers send messages to exchanges
- Exchanges distribute messages
 - One queue
 - Multiple queues
 - No queues
- Determined by *exchange type*
 - Default exchange is ""
- For pub-sub, we use a *fanout* exchange
 - Also *direct*, *topic*, *headers*

Publish-Subscribe

- Fanout exchange sends a message to all queues it knows of:
- In publisher and subscriber:
 - `channel.exchangeDeclare("logs", "fanout");`
- In publisher:
 - `channel.basicPublish("logs", "", null, message.getBytes());`
- In subscriber:
 - `String queueName = channel.queueDeclare().getQueue();`
 - `channel.queueBind(queueName, EXCHANGE_NAME, "");`

Publish-Subscribe




- Subscriber creates and binds a temporary queue per subscriber to the exchange
- Temp Q is:
 - non-durable
 - autodelete
 - generated queue name
- When a new subscriber connects to the exchange, it only sees newly published messages

Message Filtering

- Send messages to queues based on attributes/content
 - Associate Binding key with a queue in subscriber
- Use a direct exchange
 - message goes to the queues whose binding key exactly matches the message routing key
 - Message discarded if no match
- A queue can specify more than one binding key

Message Filtering



Message Filtering

- In publisher

```
channel.exchangeDeclare(EXCHANGE_NAME, "direct");  
channel.basicPublish(EXCHANGE_NAME, "black", null,  
message.getBytes());
```

- In one subscriber


```
channel.exchangeDeclare(EXCHANGE_NAME, "direct");  
String queueName = channel.queueDeclare().getQueue();  
channel.queueBind(queueName, EXCHANGE_NAME, "black");  
channel.queueBind(queueName, EXCHANGE_NAME, "green");
```

- In other subscriber

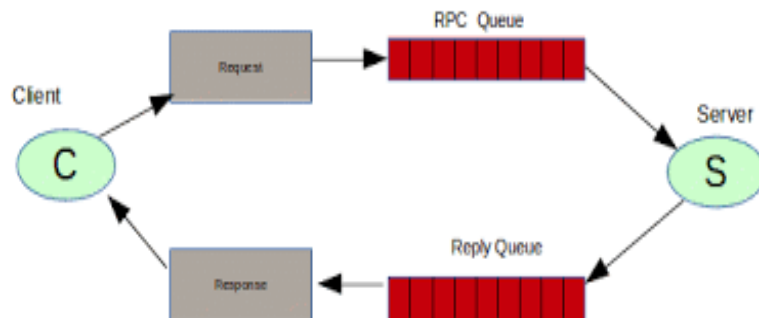
```
channel.exchangeDeclare(EXCHANGE_NAME, "direct");  
String queueName = channel.queueDeclare().getQueue();  
channel.queueBind(queueName, EXCHANGE_NAME, "orange");
```

Multi-criteria Filtering

- Filtering based on multiple criteria
- Topic exchanges
 - Khoury.staff.announce
 - Khoury.student.announce
 - Khoury.student.party
 - #.announce
 - Khoury.*.announce
 - Khoury.student.*



RabbitMQ RPC – Send and Receive



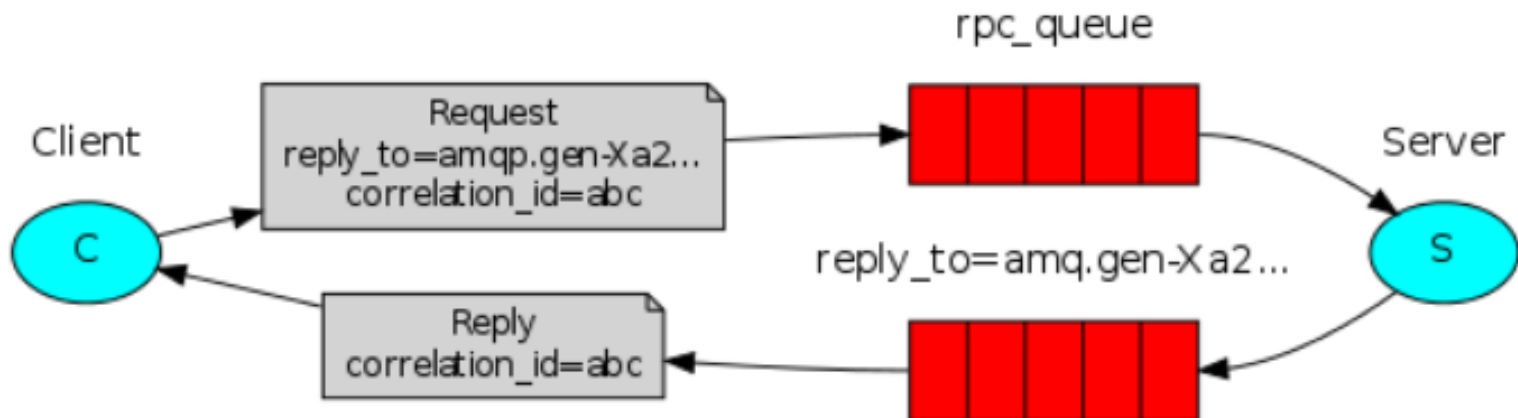
- Emulate a synchronous call with queues
- Client sends request via a queue and specifies a *callback queue*
- Server processes the message and sends result to callback queue

Correlation IDs

- How does client know which request a response is related to?
 - All in same response queue
 - Might not be in the same order as sent
- Solution is a Correlation ID
 - Unique value for every request
 - Sent in response message
 - Enable matching response to request



Correlating Requests and Responses



RabbitMQ RPC

- Client sends message with two properties:
 - replyTo: anonymous exclusive queue created just for the request,
 - correlationId: unique value for every request.
- The request is sent to an rpc_queue queue.
- The RPC worker (aka: server) waits for requests on rpc_queue. On Arrival:
 - processes request
 - sends a message with the result back to the Client, using the queue from the replyTo field, includes correlationId
- Client waits on the reply queue. On arrival:
 - checks the correlationId property matches the value from the request
 - Processes the response

RabbitMQ RPC Client

- [Server](#) and [client](#) examples – check ‘em out 😊

// construct message with necessary properties

```
String replyQueueName = channel.queueDeclare().getQueue();
```

```
    AMQP.BasicProperties props = new AMQP.BasicProperties
```

```
        .Builder()
```

```
        .correlationId(correlationId)
```

```
        .replyTo(replyQueueName)
```

```
        .build();
```

```
channel.basicPublish("", requestQueueName, props, message.getBytes("UTF-8"));
```

RabbitMQ RPC Server


```
// on message arrival
```

```
AMQP.BasicProperties replyProps = new AMQP.BasicProperties  
    .Builder()  
    .correlationId(delivery.getProperties().getCorrelationId())  
    .build();
```

```
// process the message and create response
```

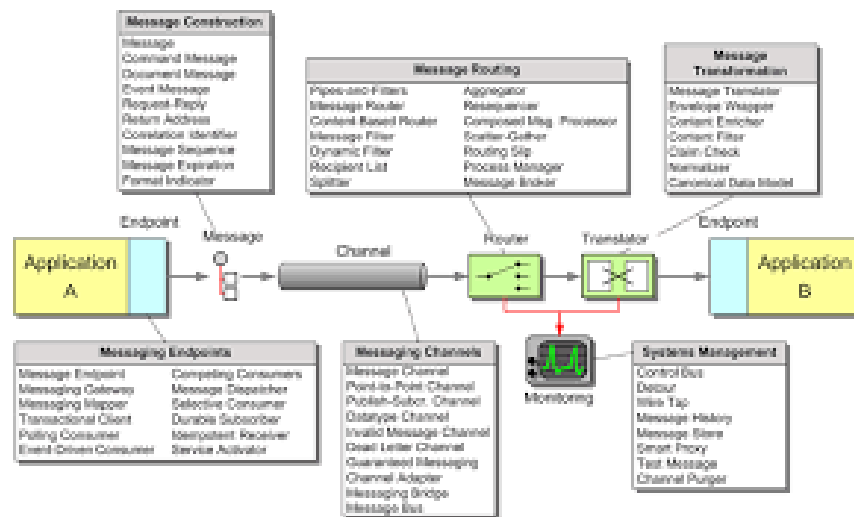
```
channel.basicPublish("", delivery.getProperties().getReplyTo(), replyProps,  
    response.getBytes("UTF-8"));  
channel.basicAck(delivery.getEnvelope().getDeliveryTag(), false);
```

Lab




Messaging Patterns

Messaging Patterns Catalog



- Catalog of ~65 patterns at [Enterprise Integration Patterns](http://EnterpriseIntegrationPatterns.com)
- Comprehensive
- Some of these important for scalability



Competing Consumers Pattern

Advantages

- **Scalability:** consumers can be increased or decreased on the fly.
- **Availability:** If consumers are unresponsive or overloaded, the queue can still store messages
- **Guaranteed Delivery:** At least once
- **Failover:** If a consumer fails mid-task, the message is returned to the queue, to be picked up by another consumer.




Implications



Message Ordering: no guarantee of the order in which messages are processed



Poison Messages: how do we handle malformed messages that cause crashes/exceptions?



Idempotence: consumers must implement idempotent processing

At least once processing

- In RabbitMQ
 - A consuming application should acknowledge a messages when all processing completed.
 - Once ack received, the broker is free to mark the message for deletion from the queue
- If an acknowledgement is not received
 - RMQ [detects connection/node failures](#)
 - Delivers message to another worker
 - Marks message with *redelivered* flag
- Ergo, processing must be idempotent

Message Invisibility Window

- Messages hidden in queue until processing confirmed
- In some systems, (eg RMQ), messages deleted by broker when acknowledgement received
- In others, (eg AWS SQS), consumer must delete message within invisibility window time
 - Default 30 seconds
- If not deleted, message becomes visible for processing again
 - Ergo idempotence


Poison Messages

- Corrupted messages, malicious content, programming bugs -> all may cause workers to fail
 - Poison messages
- Workers continue to fail on redelivery
 - Not good!!
- Somehow must detect these and stop them from being continually delivered
 - Deliver to *dead.letter* queue
 - Set an alarm
 - Diagnose problem

Detecting Poison Messages


- In RMQ:
 - If a client dies while processing a message, message is requeued
 - Next time it is delivered a redelivered header value is set
 - Client may also send a 'reject' acknowledgement and specify that the message should not be requeued
 - Clients can send poison messages to a DLX
- In other systems (e.g, AWS SQS)
 - Message have a delivery count
 - If delivery count > n (3?), send to dead letter queue

Scaling Messaging



- Broker clustering in RMQ
- Brokers act as peers
 - Authenticate
 - replicate state
- Queues not replicated
- Clients connect to any broker and requests are forwarded to queue origin

Queue Mirroring



- RMQ enables queues to be mirrored across cluster members
- Each queue has a master copy and is mirrored to other secondaries
 - Configurable batch sizes
- Producers connect to master
 - Ensures FIFO ordering
 - Mirrors for data safety
 - Consumers can read from master or mirrors

Master Failure

- If master fails, a mirror is chosen as new master
 - Most up to date mirror
 - Any non-replicated messages from master are lost
- Upon election, new master:
 - Makes all unacknowledged messages available on queue
 - Why?
- Clients must be aware of possible message redelivery



Summary



Asynchronous systems can be built with messaging



Point to Point



Publish-Subscribe



Many messaging platforms



Messaging patterns



Clustering and mirroring