

Northeastern University - Seattle



CS6650 Building Scalable Distributed Systems
Professor Ian Gorton

Building Scalable Distributed Systems

Week 1 – Introduction to Scalable Systems

(Part 3)

What is Scalability?


Scale is being driven by ...

Around 40% of the world population has an internet connection today ([view all on a page](#)). In 1995, it was less than 1%.

The number of internet users has increased tenfold from 1999 to 2013.

The **first billion** was reached in 2005. The **second billion** in 2010. The **third billion** in 2014.

The chart and table below show the number of global internet users per year since 1993:




As well as




Data Volumes

BIG IN GROWTH, TOO.

1 exabyte (EB) = 1,000,000,000,000,000 bytes



Number of Connected Objects Expected to Reach 50bn by 2020



Penetration of connected objects in total 'things' expected to reach 2.7% in 2020 from 0.6% in 2012

Source: CCS, 2013

© 2013 Cisco and/or its affiliates. All rights reserved.

Cisco Public

2025???

<http://www.zdnet.com/article/the-internet-of-things-and-big-data-unlocking-the-power/>

Question

- How are usage statistics derived in a scalable system?
- Information must be gathered about every single request
 - Session length
 - Session origin
 - Session behavior
- In pornhub.com's case, billions of visits a year
- Lot of data
 - A big data system?



The future is about
Big Data
Big Problems
@Web Scale

Web scale



Globally accessible



(Practically) infinite user base



Constant usage (24x7)



Ability to handle surges and spikes in requests



Controllable costs

Scalability

- The ability of a system to increase or decrease capacity in response to changing demands, while continuing to satisfying service level agreements for latency and availability
- [Hyperscalability](#) – as above, while supporting exponential growth with linear costs


Just need more machines and disks?



**Major Internet
Companies
have 1m+ servers in
2013**

And Scale Linearly



Scale-Up Linearity



<http://www.datastax.com/2012/01/choosing-the-right-architecture-for-big-data-scale>

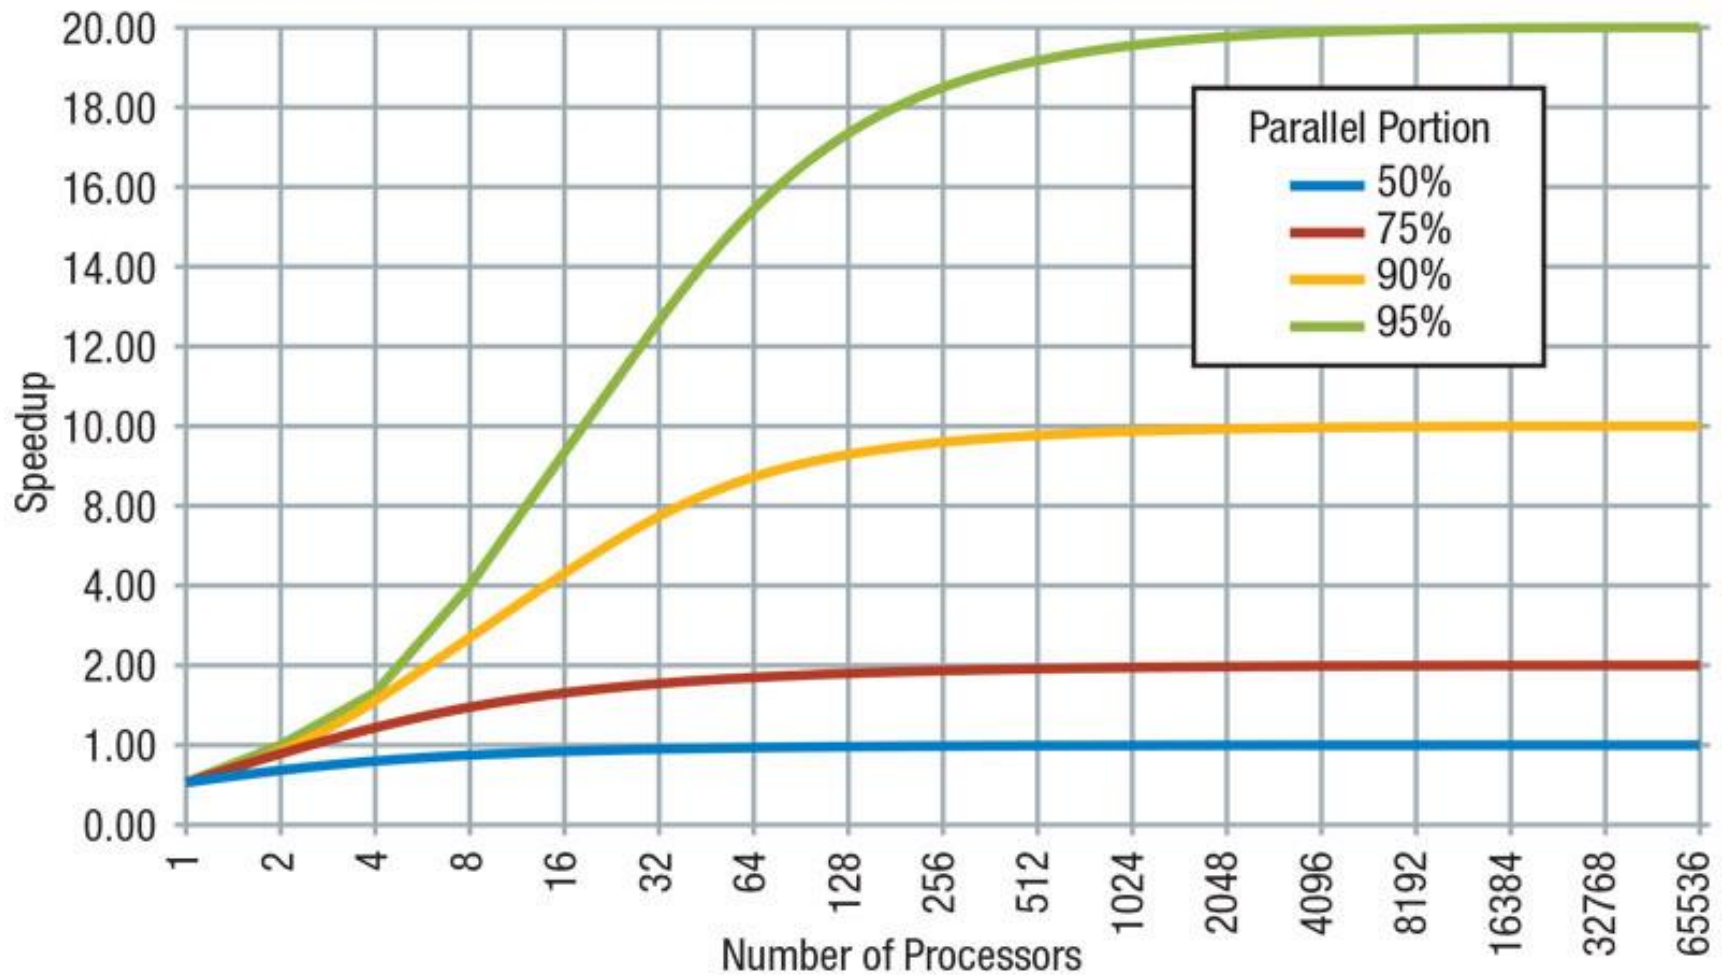


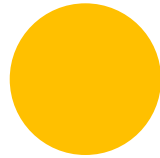
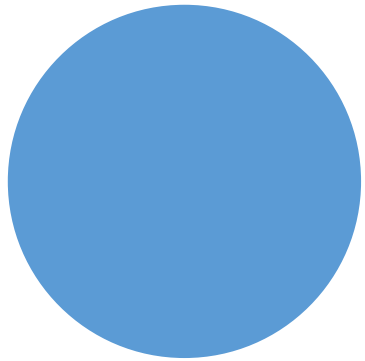
Anyone
heard of
Amhdahl's
law?



Amhdahl's Law

Amdahl's Law






Amhdahl's Law in
2025???



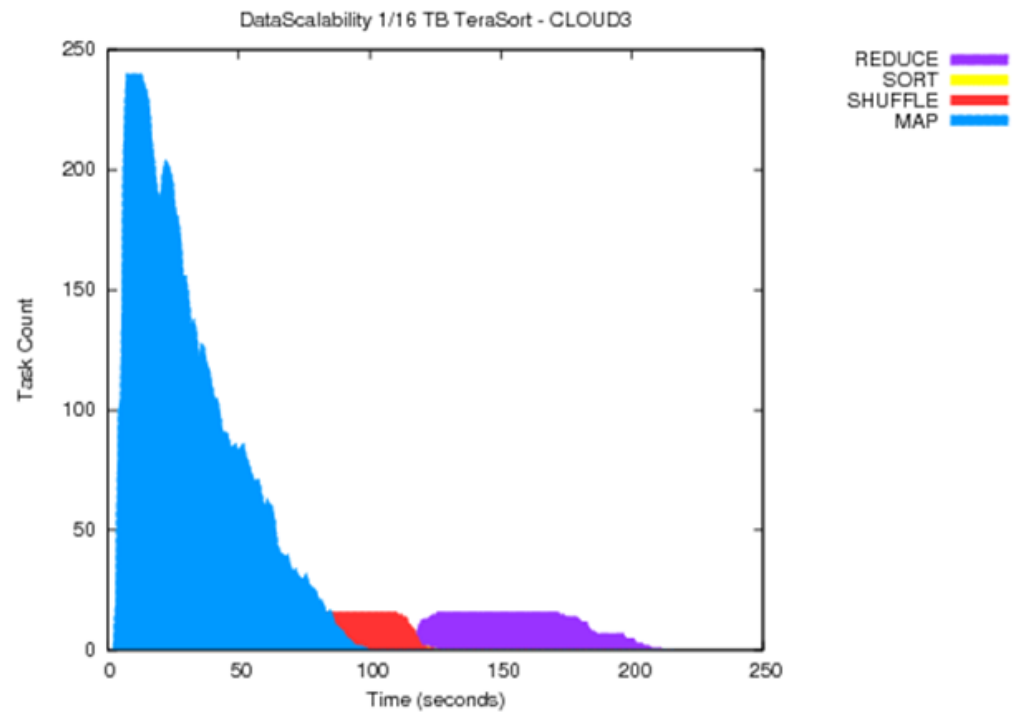



**NO
CHANGE**

Many
problems
are
easily
parallelized
though?



Data Skew???





The future is Big Problems

That are difficult to design

And difficult to build

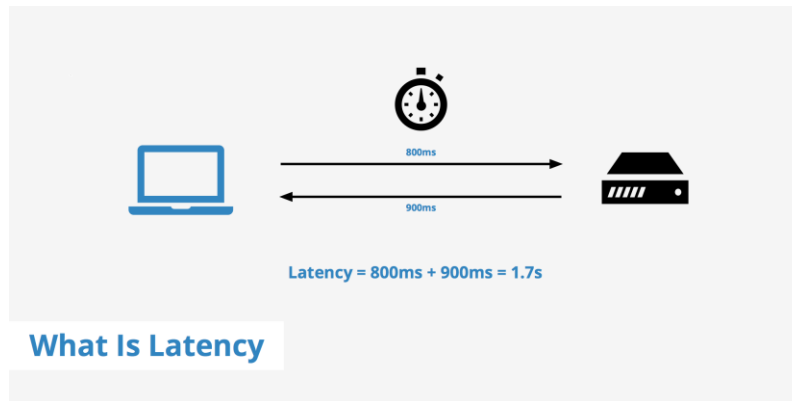
And difficult to deploy



Engineering Challenges

- High traffic:
 - Consider ad exchanges that are responding to 1,000,000 requests/second
- Huge volumes of data:
 - Often terabytes or petabytes of data
- Strict requirements for systemic properties:
 - Response time, throughput, security, availability, compute cost, ...

Response time



- Also known as latency
- Round trip time from when a request sent until a response is received

Response Time



Response time has an impact on user behavior

Many studies show a decrease in sales and customer retention as the response time increases

<http://www.martinlugton.com/page-load-speed-important-impact-site-speed-conversions-revenue/>

“Google: increasing page load time from 0.4 seconds to 0.9 seconds decreased traffic and ad revenues by 20%”



In some cases it's required by contract

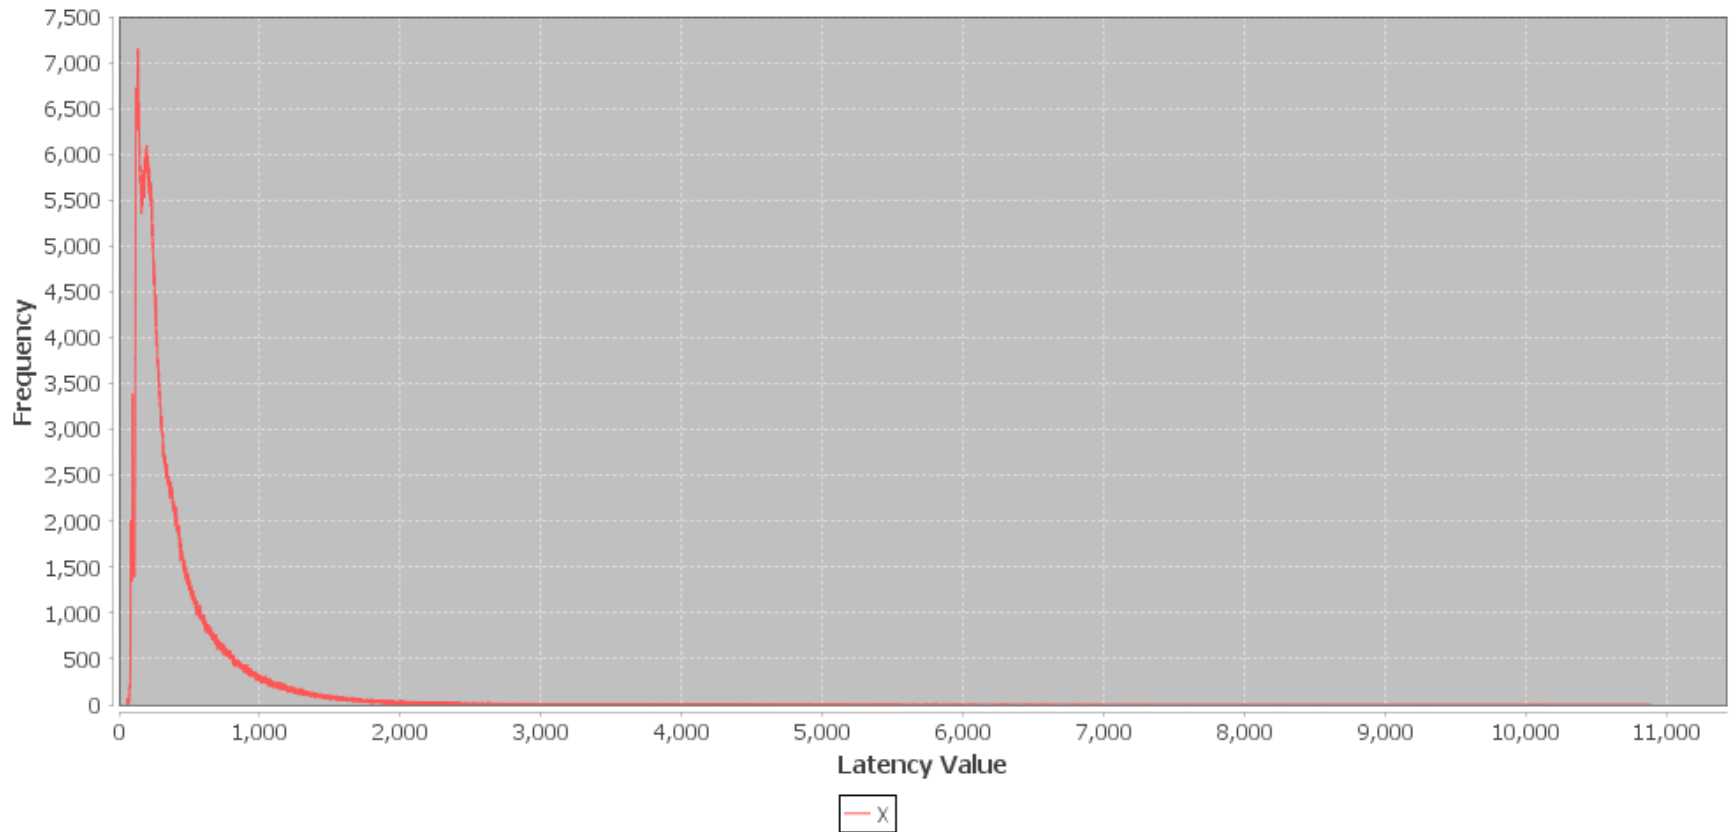
Responses to ad-exchanges for example need to be within 40 milliseconds



Website response time impacts google rankings


One factor google uses when determining ad listings

POST-latency-plot



Typical Response Time – Long Tail

Throughput



- The capacity of a system to process requests
- Typically measured as number of requests/second
- Cf latency:
 - Latency is a measure for a single request
 - Throughput is a bulk measure of many requests

Throughput



Massive throughput is needed at scale

Millions of requests per second



Before web scale systems, we just used bigger more powerful systems

Particularly for the database server




That's no longer an option at web scale




We need to be able to efficiently scale out

In other words be able to split the load across multiple nodes

Throughput with fixed capacity



Scale Out



- **Scale** horizontally (or **scale out**) means adding more nodes to a system
- Increases capacity and hence throughput
- 100s to 1000s of nodes needed at scale
- Scale in/down equally important
 - Why???

Availability



**The percentage of time a system
is accessible to user**



At web scale

Global customer base
High availability required



How high?

99%
99.999%
99.999999%

Availability

- If a system has 1 component
 - It is not available if the element (software/hardware) fails
 - Single Point of Failure (SPoF)
 - So need very highly reliable component
- To achieve scalability, we need to scale out
 - 100s and 1000s of components
 - Do they all need to be always available?
- At scale *everything* breaks ...
 - 3% of disks fail in a data center each year
- Systems need to be able to deal with regular component failures
 - In other words, *fault tolerant*
 - ***Built to handle problems as failure is expected***

Managing Tradeoffs

- There is not one way to build such systems
- Requirements and context are going to differ from one system to another
- Optimizing on one system property is inevitably going to compromise another concern, eg
 - Throughput versus cost
 - Performance versus availability
- As an engineer you need to be able to understand:
 - the concerns and context of a given system
 - the nature of the tradeoff that results from a set of decisions
 - And use metrics to guide your designs

Summary



This course examines the trade-offs needed to build a scalable system



We do this though studying scalable distributed system principals and technologies that enable them



We build a series of assignments which give practical insights into building scalable systems



Next – course outline